

Topic 2: Classification

Eric B. Laber

Department of Statistical Science, Duke University

Statistics 561



On choices

I'd rather be rich than stupid.
—Pierre-Joseph Proudhon



On having to cover classification in two lectures

I needed to think last night. So I galloped into a wooded glen, and after punch dancing out my rage and suffering an extremely long and very painful fall, I realized what has to be done.

—Rod




Warm-up (5 minutes)

- ▶ Explain to your group
 - ▶ What is classification? What are some canonical examples?
 - ▶ What is the separability problem?
 - ▶ What is zero-one loss?
- ▶ True or false
 - ▶ Regression is a misnomer in 'logistic regression'
 - ▶ Classification is in the "inner loop" of many RL algorithms
 - ▶ McDonald's is the most prevalent fast-food chain in NC

Classification: quick overview

- ▶ Input-output pairs where output is one of finitely many categories
 - ▶ High-risk v low-risk for complications in surgery
 - ▶ Spam v not spam
 - ▶ Handwritten digit recognition
- ▶ Example from sex-trafficking classification problem



Setup: binary classification

- ▶ Observe $\{(\mathbf{X}_i, Y_i)\}_{i=1}^n$ comprising i.i.d. draws from P
 - ▶ Inputs: $\mathbf{X} \in \mathbb{R}^p$
 - ▶ Outputs: $Y \in \{-1, 1\}$, aka, label
- ▶ Classifier $c : \mathbb{R}^p \rightarrow \{-1, 1\}$ so that $c(\mathbf{x})$ is the predicted label at input $\mathbf{X} = \mathbf{x}$

0-1 loss

- Natural measure of classification performance is 0-1 loss

$$\ell_0(\mathbf{x}, y; c) \triangleq 1_{y \neq c(\mathbf{x})}$$

so that the expected loss (risk) is

$$\tau(c) \triangleq P\ell_0(\mathbf{X}, Y, c) = P1_{Y \neq c(\mathbf{X})} = P\{Y \neq c(\mathbf{X})\}$$

- E.g., linear classifier $c(\mathbf{x}; \beta) = \text{sign}(\mathbf{X}^\top \beta)$ which has loss

$$\tau(\beta) \triangleq P1_{Y \neq \text{sign}(\mathbf{X}^\top \beta)} = P1_{Y\mathbf{X}^\top \beta < 0}$$

Bayes classifier

- Let \mathcal{C} be class of all (msbl) maps from \mathbb{R}^p into $\{-1, 1\}$ and

$$c^{\text{opt}} = \arg \min_{c \in \mathcal{C}} \tau(c) = \arg \min_{c \in \mathcal{C}} P\{Y \neq c(\mathbf{X})\}$$

then

$$c^{\text{opt}}(\mathbf{x}) = \begin{cases} 1 & \text{if } P(Y = 1|\mathbf{X} = \mathbf{x}) \geq 1/2 \\ -1 & \text{otherwise} \end{cases}$$

i.e., $c^{\text{opt}}(\mathbf{x}) = \text{sign}\{2q(\mathbf{x}) - 1\}$ w/ $q(\mathbf{x}) \triangleq P(Y = 1|\mathbf{X} = \mathbf{x})$

Blank page for notes



Blank page for notes



Probabilistic classifiers

- ▶ Natural approach to classification is to estimate construct an estimator $\hat{q}_n(\mathbf{x})$ of $q(\mathbf{x}) = P(Y = 1|\mathbf{X} = \mathbf{x})$
- ▶ Logistic regression posits a model of the form

$$q(\mathbf{x}; \boldsymbol{\beta}^*) = \text{expit}(\mathbf{x}^\top \boldsymbol{\beta}^*) = \frac{\exp(\mathbf{x}^\top \boldsymbol{\beta}^*)}{1 + \exp(\mathbf{x}^\top \boldsymbol{\beta}^*)}$$

Derive log-likelihood for logistic regression



Blank page for notes



In-class exercise (5 min)

- ▶ In your group derive a stochastic gradient descent algorithm for logistic regression: then on to `logistic_regression.R`

Separability problem

- ▶ If classes linearly separable estimates diverge
- ▶ Use ridge with logistic regression¹
- ▶ Penalized negative log-likelihood (0-1 coding for simplicity)

$$\ell_n(\beta) = -\mathbb{P}_n[\mathbf{X}^\top \beta Y + \{1 - q(\mathbf{X}; \beta)\}] + \lambda \|\beta\|^2$$

differentiate wrt β to obtain

$$-\mathbb{P}_n \mathbf{X} \{Y - q(\mathbf{X}; \beta)\} + 2\lambda \beta$$

¹That's right, I'm stating this without qualification of any kind. What's it to you? That's what I thought.

Large-margin classifiers

- ▶ Decision boundary $\{\mathbf{x} : q(\mathbf{x}) = 1/2\}$
- ▶ Don't need $P(Y|\mathbf{X} = \mathbf{x})$ just $\text{sign}\{2q(\mathbf{x}) - 1\}$, i.e., we don't need to know the boundary just which side we're on
- ▶ Idea! Construct loss functions that penalize distance from correct side of boundary

Surrogate loss functions: laber draws a picture

Surrogate loss functions

- ▶ Common surrogate loss functions include
 - ▶ Squared error loss: $\ell(\mathbf{x}, y; \beta) = (1 - y\mathbf{x}^\top \beta)^2$
 - ▶ Exponential loss: $\ell(\mathbf{x}, y; \beta) = \exp(-y\mathbf{x}^\top \beta)$
 - ▶ Hinge loss: $\ell(\mathbf{x}, y; \beta) = (1 - y\mathbf{x}^\top \beta)_+$
 - ▶ Logistic loss: $\ell(\mathbf{x}, y; \beta) = \log \{1 + \exp(-y\mathbf{x}^\top \beta)\}$
- ▶ Can also consider penalized versions of these loss functions

Nice and relaxing derivation

- Show $\hat{\beta}_n = \arg \min_{\beta} \mathbb{P}_n \log \{1 + \exp(-Y\mathbf{X}^\top \beta)\}$ recovers the logistic regression estimator

Blank page for notes



Statistical inference

► Convexity buys you a lot, many derivations much simpler

► High-level overview of assumptions

(A1) $\ell(\mathbf{x}, y; \beta)$ is convex in β for each (\mathbf{x}, y)

(A2) $Q(\beta) = P\ell(\mathbf{X}, Y; \beta)$ exists and is finite for all β

(A3) $\beta^* = \arg \min Q(\beta)$ exists and is unique

(A4) $Q(\beta)$ is twice continuously differentiable in a nbrhd of β^* and $H = \nabla^2 Q(\beta^*)$ is positive definite

Theorem

Assume (A1)-(A4) and let $\Omega = P\nabla\ell(\mathbf{X}, Y; \beta^*)\nabla\ell(\mathbf{X}, Y; \beta^*)^\top$ then

$$\sqrt{n}(\hat{\beta}_n - \beta^*) \rightsquigarrow \text{Normal}(0, H^{-1}\Omega H^{-1}).$$



Statistical inference cont'd

- ▶ Asymptotic normality \Rightarrow we can use standard methods of inference for β^* as in the regression case, e.g., Wald-type intervals etc.
- ▶ Unlike regression case, measures of performance are not well-behaved
- ▶ Recall our measures of performance
 - ▶ Population-level error: $R \triangleq P1_{Y\mathbf{X}^\top\beta^* < 0}$
 - ▶ Conditional error: $C(\hat{\beta}_n) \triangleq P1_{Y\mathbf{X}^\top\hat{\beta}_n < 0}$
 - ▶ Average error: $A_n \triangleq \mathbb{E}C(\hat{\beta}_n)$

Comparing measures of performance

- ▶ Fact: the three measures of performance need not coincide even in infinite samples

Blank page for notes



Blank page for notes



Thank you.

`eric.laber@duke.edu`

`laber-labs.com`

